

Unsupervised Point-Cloud Reconstruction

Luca Barco s276072@studenti.polito.it
Stefano Bergia s276124@studenti.polito.it
Daniela De Angelis s277493@studenti.polito.it


Politecnico di Torino
Machine Learning and Artificial Intelligence
A.Y. 2020/2021



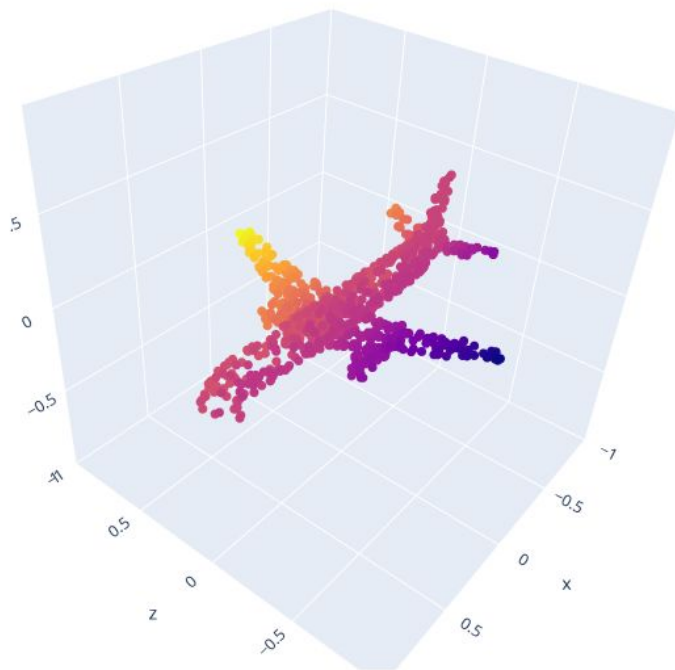
Image: polito.it

What is a **Point-Cloud**?

- Set of points
- Unordered and unstructured
- Represents an object
- Close to sensor data (e.g. LiDAR Camera)



Cannot directly use traditional CNN Models




How deep learning methods deal with **Point-Clouds**?

To deal with Point-Clouds, deep learning methods must guarantee:

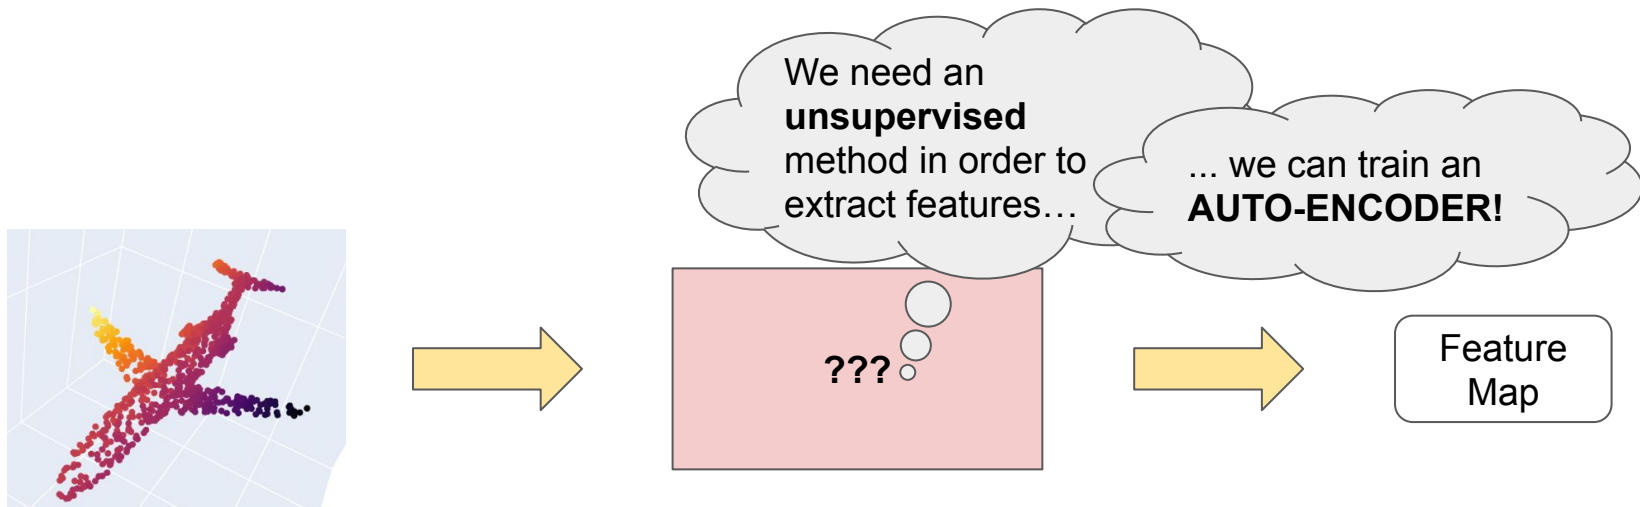
- **Permutation Invariance**: invariant to $N!$ permutations of the input.
- **Geometric Transformations Invariance**: Rigid transformation (e.g. rotation) applied to the input should not alter the performed task results.

Build a Point-Cloud Auto-Encoder

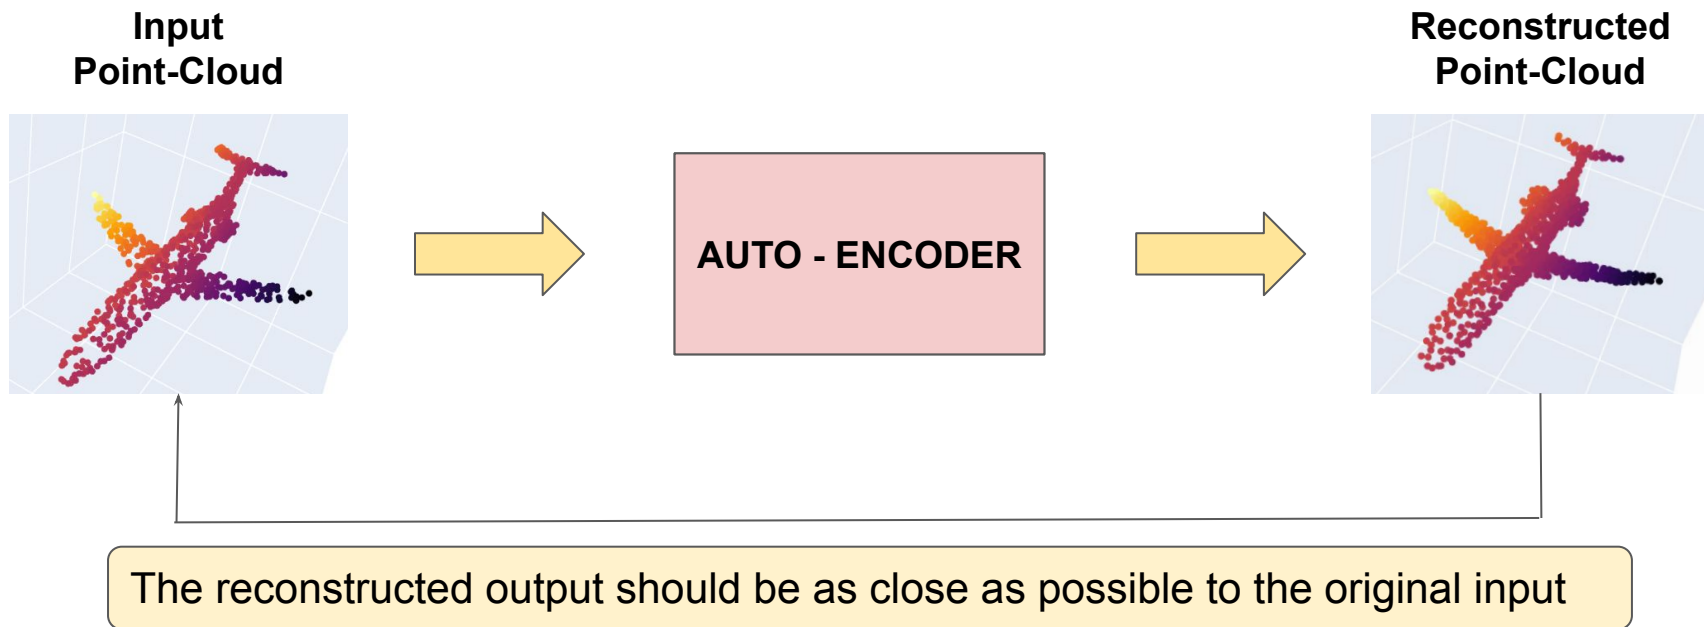
How to **extract** information from a Point-Cloud?



How to **extract information** from a Point-Cloud?



Our problem: train a Point-Cloud Auto-Encoder



Starting point: PointNet and DGCNN

PointNet^[1] and **DGCNN^[2]** are **supervised** deep learning method for Point-Clouds classification and segmentation.



They extract the main features of a Point-Cloud

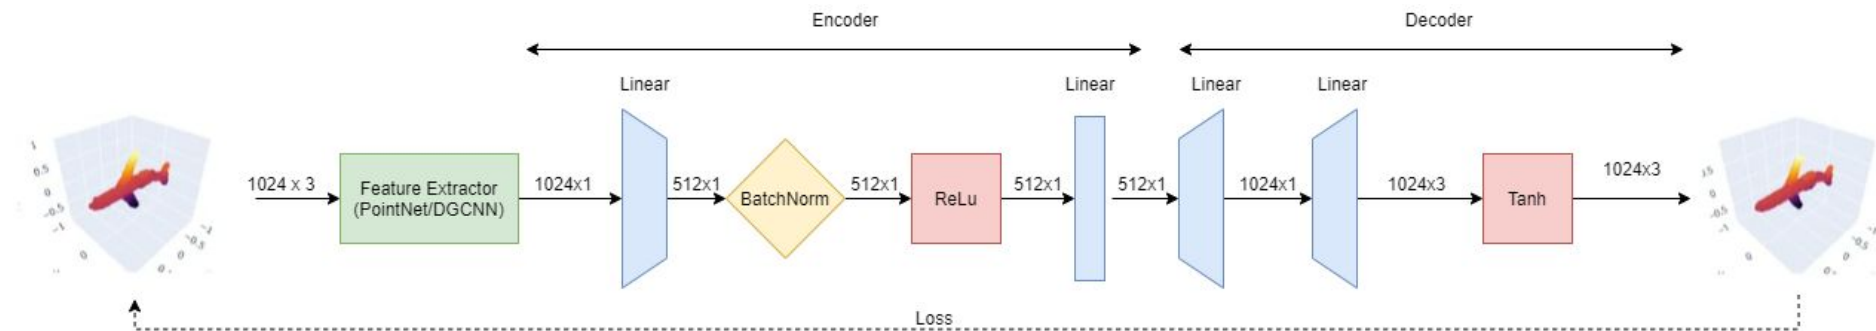


Use their features extractors in an **unsupervised** Auto-Encoder

[1] Qi et Al., PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation, in CVPR 2017

[2] Y. Wang, Y. Sun, Z. Liu, S. E. Sarma, M. M. Bronstein, and J. M. Solomon. Dynamic graph CNN for learning on point clouds. CoRR, abs/1801.07829, 2018.

Auto-Encoder's Architecture



Main components:

- **Feature Extractor** (based on PointNet^[1] or DGCNN^[2])
- **Encoder**
- **Fully Connected Decoder**
- Loss function based on **Chamfer Distance**

[1] Qi et al., PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation, in CVPR 2017

[2] Y. Wang, Y. Sun, Z. Liu, S. E. Sarma, M. M. Bronstein, and J. M. Solomon. Dynamic graph CNN for learning on point clouds. CoRR, abs/1801.07829, 2018.

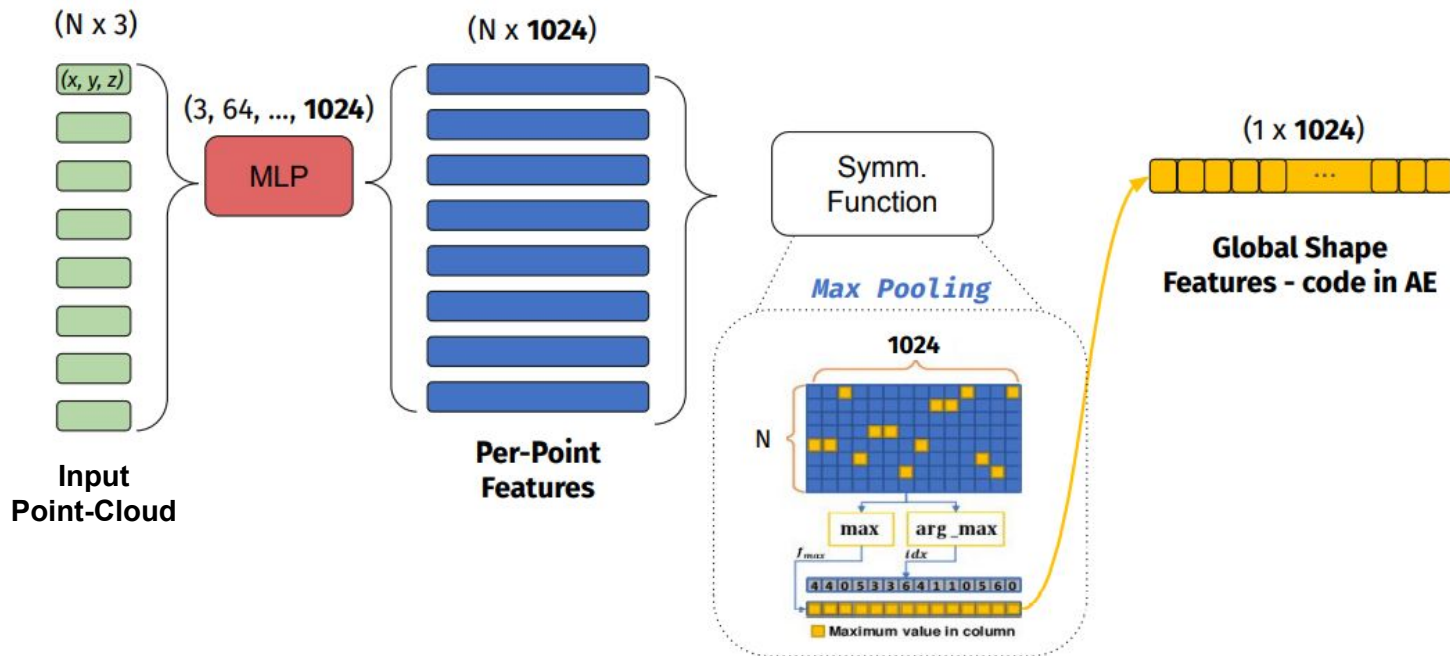
Auto-Encoder's Architecture: PointNet^[1] Feature Extractor

PointNet^[1]: deep learning method for Point-Clouds classification and segmentation.

- **Permutation Invariance** : guaranteed using a symmetric function to aggregate information from each of the points (e.g. Max Pooling).
- **Geometric Transformations Invariance**: guaranteed using Transformer Networks (T-Net) in euclidean space.
- It aggregates **global features** into a unique vector

Auto-Encoder's Architecture: PointNet^[1] Feature Extractor

PointNet^[1] Architecture: Features extraction



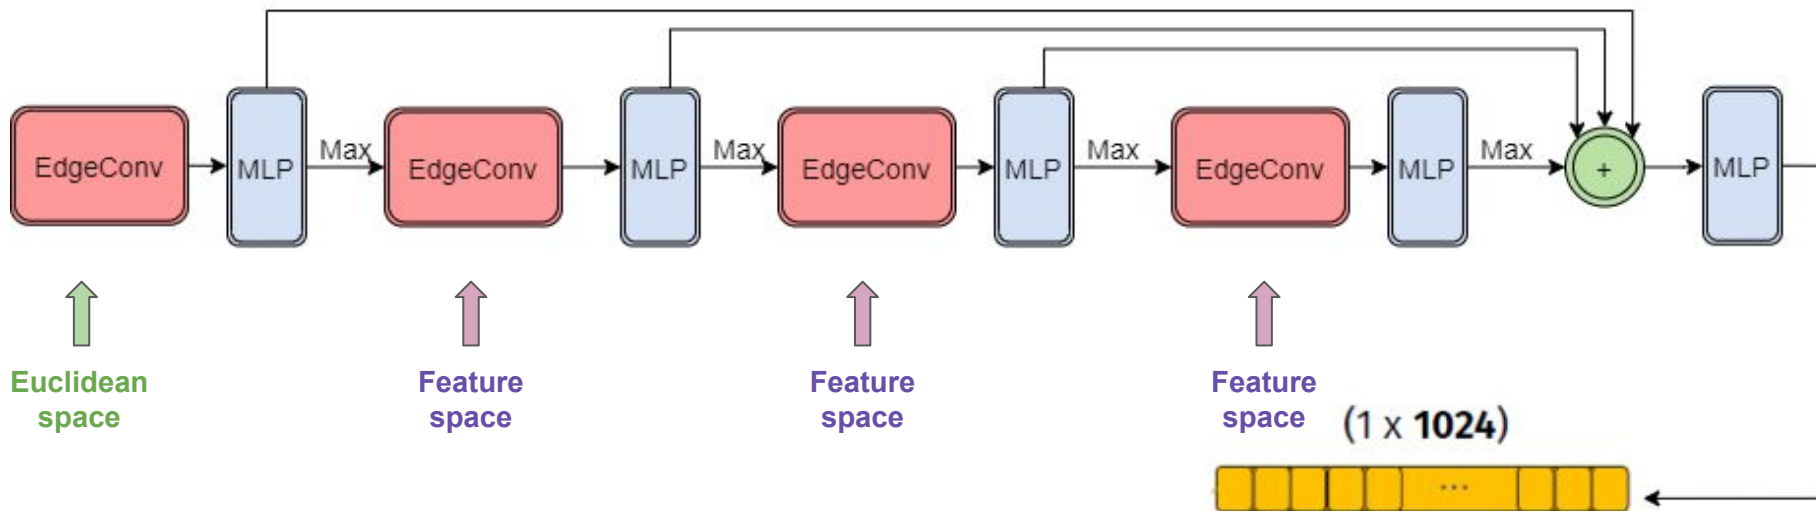
Auto-Encoder's Architecture: DGCNN^[2] Feature Extractor

DGCNN^[2]: deep learning method for Point-Clouds classification and segmentation based on computation of direct graph representing local structures.

- **Permutation Invariance** : guaranteed using a symmetric function (e.g. Max Pooling) to aggregate information at each layer.
- **Geometric Transformations Invariance**: guaranteed using the neighbourhood graph.
- It aggregates both **global and local** features into a unique vector

Auto-Encoder's Architecture: DGCNN^[2] Feature Extractor

- **Dynamic Graph:** the graph is updated at each layer
- **EdgeConv:** convolutional layer working on the dynamic graph



Auto-Encoder's Architecture: DGCNN^[2] Graph Computation

How is the graph computed?


- Using an **Edge function**: asymmetric function that considers each point x_i and its neighbourhood ($x_j - x_i$)

$$h_{\Theta}(x_i, x_j) = \bar{h}_{\Theta}(x_i, x_j - x_i).$$

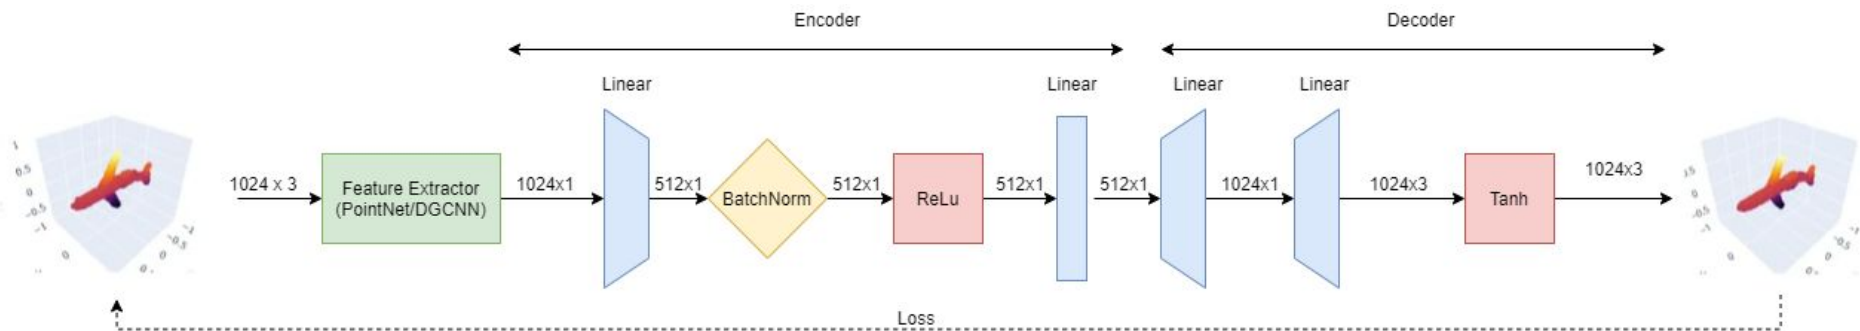
Θ : Learnable Parameters

Global Shape

Local Neighborhood Structure



Auto-Encoder's Architecture: Encoder and Decoder



- **Encoder:** Further refinement of features of the input Point-Cloud
- **Decoder:** Fully-Connected layers to reconstruct the input Point-Cloud

Auto-Encoder's Architecture: Loss Function

Distance metric: **Chamfer Distance**^[3]

Given two point sets $P_1, P_2 \rightarrow$

$$d_{CD}(P_1, P_2) = \sum_{x \in P_1} \min_{y \in P_2} \|x - y\|_2^2 + \sum_{y \in P_2} \min_{x \in P_1} \|x - y\|_2^2$$

$P_1 \rightarrow P_2$ $P_2 \rightarrow P_1$

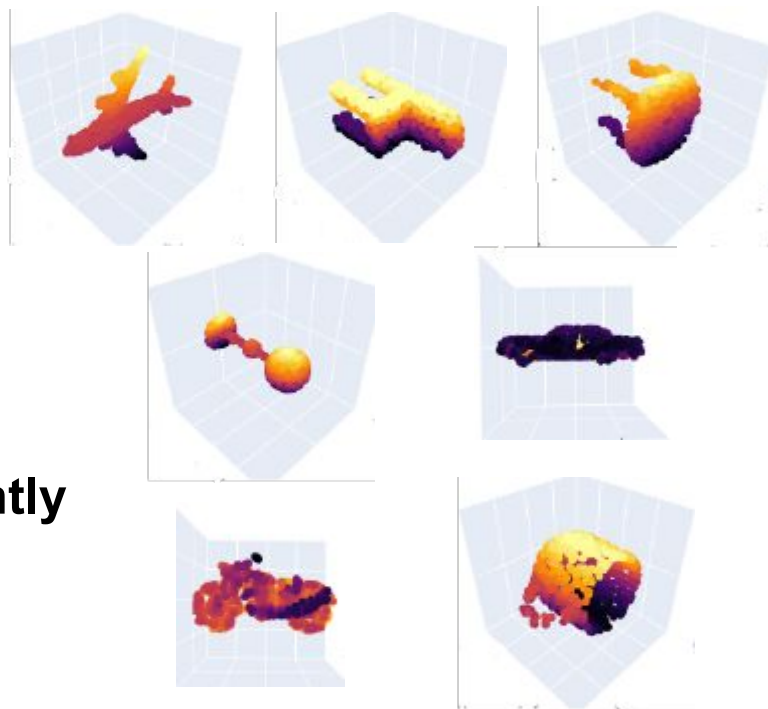
- “pseudo-symmetrical” behaviour:
 - For each point in P_1 , compute the squared distance with respect to the nearest point in P_2
 - The same for each point in P_2
 - Sum these two components

Loss Function: average of the Chamfer Distance over N points of Point-clouds P_1 and P_2 multiplied by a w factor (in our work $w=100$)

$$L(P_1, P_2) = \frac{w}{2N} \cdot d_{CD}(P_1, P_2)$$

Auto-Encoder: Training

- Training Dataset: **ShapeNet**^[4]
- Training on seven classes of data:
Airplane, Chair, Table, Lamp,
Car, Motorbike, Mug.
- **Two scenarios:**
 - **Training on all seven classes jointly**
 - **Training on single classes**



Auto-Encoder: Training on all 7-classes vs single

- Two Architectures:
 - **PNet_AE_512** : PointNet based AE, Encoder lower level size: 512
 - **DGCNN_AE_512**: DGCNN based AE, Encoder lower level size: 512

**Training on all seven classes jointly:
Testing loss results**

Category	PNet_AE_512	DGCNN_AE_512
Airplane	0.100	0.105
Chair	0.191	0.190
Table	0.207	0.219
Lamp	0.301	0.272
Car	0.245	0.256
Motorbike	0.186	0.193
Mug	0.381	0.364
Avg	<i>0.230</i>	<i>0.228</i>


**Training on single classes:
Testing loss results**

Category	PtNet_AE_512	DGCNN_AE_512
Airplane	0,113	0,099
Chair	0,221	0,197
Table	0,227	0,228
Lamp	0,387	0,499
Car	0,294	0,254
Motorbike	0,225	0,236
Mug	0,582	0,585
Avg	<i>0,293</i>	<i>0,300</i>


Auto-Encoder: Training on all 7-classes vs single




Regardless of the architecture, results look better when training on all 7-classes jointly!



Original




Training on single classes



Training on 7-classes jointly

Point-Cloud contains uncommon details



Loss of details: the model learns an *average shape*



Features learned on one category can be used to reconstruct others

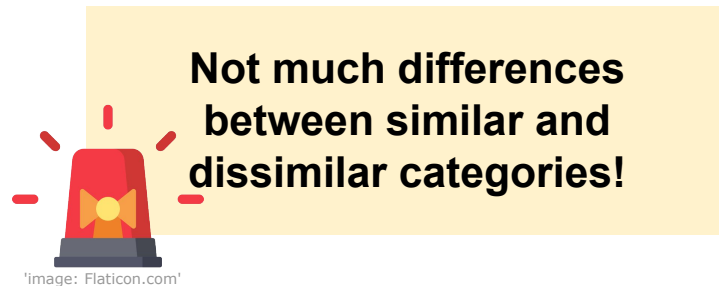
Auto-Encoder: Testing on novel categories

Testing the models trained on all seven classes jointly on novel unseen categories:

- **Similar:** *Basket, Bicycle, Bowl, Helmet, Microphone, Rifle, Watercraft*
- **Dissimilar:** *Bookshelf, Bottle, clock, Microwave, Pianoforte, Telephone*



Testing on novel categories: Testing loss results

	Category	PNet_AE_512	DGCNN_AE_512
Similar	Basket	0.711	0.606
	Bicycle	0.408	0.399
	Bowl	1.072	0.747
	Helmet	0.902	0.732
	Microphone	1.635	0.694
	Rifle	0.201	0.197
	Watercraft	0.261	0.259
	Avg	0.741	0.516
Dissimilar	Bookshelf	0.576	0.551
	Bottle	0.330	0.307
	Clock	0.702	0.562
	Microwave	0.494	0.517
	Pianoforte	0.732	0.631
	Telephone	0.494	0.424
	Avg	0.584	0.499





Auto-Encoder: Examples of outputs

Similar:



Uses Motorbike features!

Dissimilar:




Uses Car Features!

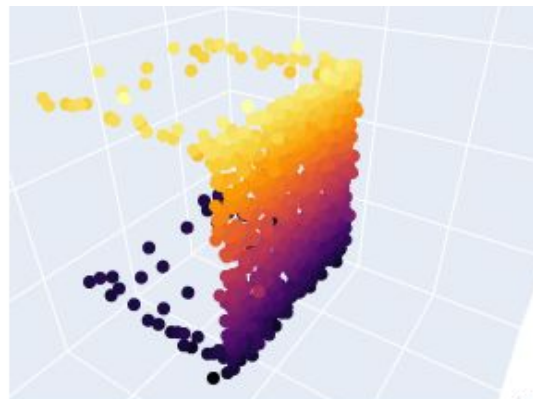
Auto-Encoder Loss Function issue: local density

Problem: it doesn't preserve local density \longrightarrow Loss of finer details


Input Point-Cloud




Reconstructed Point-Cloud



Lower densities for
legs and bars, higher
density for the plane



Errors in low density
regions are penalized
much more



Few points are
assigned to low
density regions

Auto-Encoder: Pros, cons and improvements

PROS

- ✓ Good at reconstructing general shape of the object
- ✓ Good at “transfer” local features learned for an object to reconstruct another one
- ✓ Small embedding size

CONS

- ✗ Bad at reconstructing finer details
- ✗ Bad at preserving local density



IMPROVEMENTS



Adjust the loss function to preserve density and better reconstruct finer details


Point-Cloud Completion: Complete a cropped Point-Cloud

How to **complete** a cropped Point-Cloud?



Starting point: PF-Net^[4]

PF-Net^[4] is a deep learning method to reconstruct 3-D objects considering only the missing part

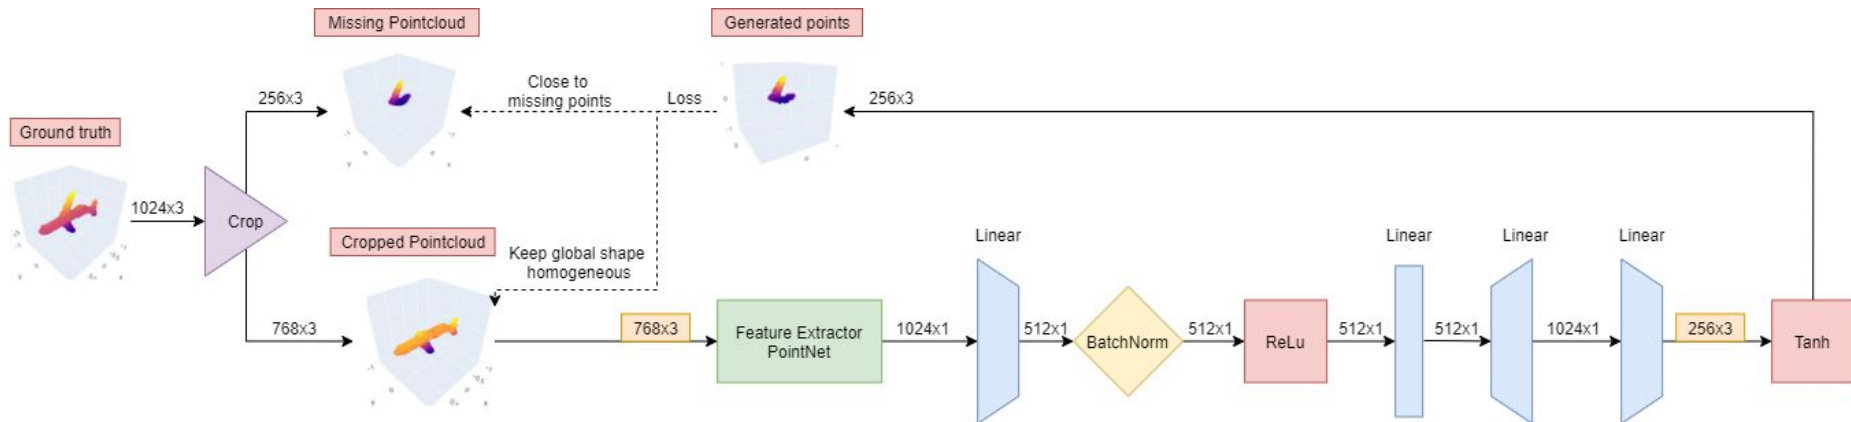


Use their idea with **our Auto-Encoder**

[4] Z. Huang, Y. Yu, J. Xu, F. Ni, and X. Le. Pf-net: Point fractal network for 3d point cloud completion, 2020.

How to **complete** a cropped Point-Cloud?

Our proposal



Key concepts:


- **Point-Cloud cropping** → reconstruct only the Missing Part
- **PointNet based Auto-Encoder**
 - **input:** Cropped Point-Cloud
 - **output:** Reconstructed Missing Part Point-Cloud
- **Final result:** Reconstructed Missing Part + Cropped → **complete object Point-Cloud**
- **Objective:** good reconstruction of Missing Part + keep global object shape homogeneous

Point-cloud completion: Point-Cloud cropping


- Set of **viewpoints**

- **For each one**

Remove the 256 nearest points



- 1) **Cropped Point-Cloud**
- 2) **Missing Part Point-Cloud**



Point-cloud completion: Loss function

- \mathbf{P}_{gt} : Ground Truth Point-Cloud
- \mathbf{P}_m : Missing-Part Point-Cloud
- \mathbf{P}_c : Cropped Point-Cloud
- \mathbf{P}_{rm} : Reconstructed Missing-Part Point-Cloud
- $\mathbf{P}_r = \mathbf{P}_c + \mathbf{P}_{rm}$: Reconstructed Point-Cloud given by concatenation of \mathbf{P}_{rm} and \mathbf{P}_c .
- $L(\mathbf{P}_1; \mathbf{P}_2)$ as the mean Chamfer Distance between Point-Clouds \mathbf{P}_1 and \mathbf{P}_2 .
- r : weighting parameters to balance the two terms

→ **Matching Loss function** $L_{PCC} = L(P_{rm}, P_m) + r \cdot (L(P_r, P_{gt}))$



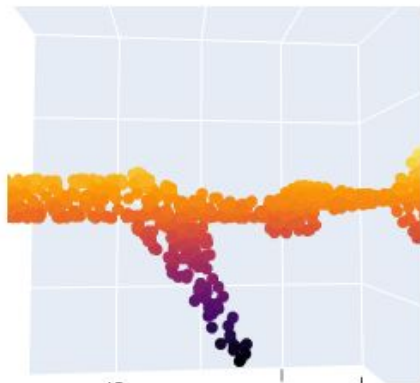
good reconstruction of missing part
(also called **Vanilla Loss Function**)



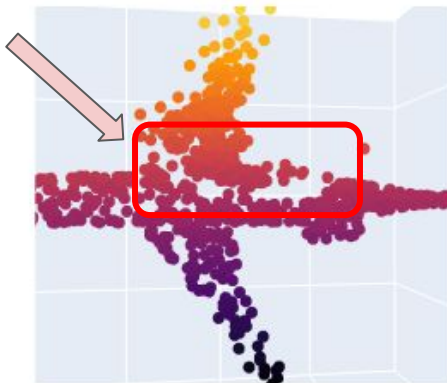
keep global shape
homogeneous

Point-cloud completion: Vanilla vs Matching loss


Cropped



Vanilla loss:
Decoded loss=0.208



Matching loss:
Decoded loss=0.221



Point-cloud completion: Training

- Training on all seven classes jointly
 - *Airplane, Chair, Table, Lamp, Car, Motorbike, Mug.*

	Category	PNet_AE_512	
		Vanilla Loss Model	Matching Loss Model
Known	Airplane	0.103	0.102
	Chair	0.181	0.167
	Table	0.180	0.191
	Lamp	0.592	0.639
	Car	0.181	0.249
	Motorbike	0.130	0.372
	Mug	0.283	0.422
	Avg	0.236	0.306

Table . Testing Loss results for Point-Cloud Completion (training on all 7 known classes jointly using Vanilla Loss and Matching Loss). Testing Loss: mean Chamfer Distance only on the missing part.

Outputs for Matching Loss model



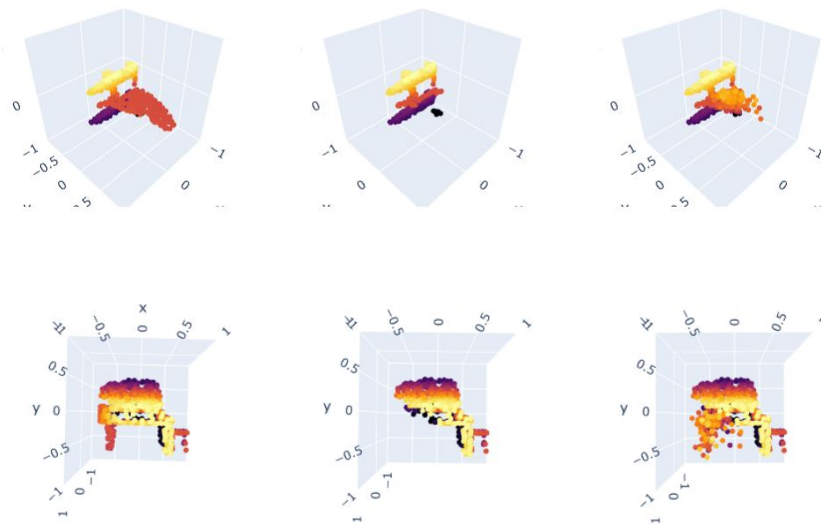
Wings
symmetry?

Point-cloud completion: Testing on novel categories

	Category	PNet-AE_512	
		Vanilla Loss Model	Matching Loss Model
Similar	Basket	0.264	0.249
	Bicycle	0.155	0.159
	Bowl	0.474	0.416
	Helmet	0.287	0.287
	Microphone	2.426	2.168
	Rifle	0.152	0.141
	Watercraft	0.087	0.094
	Avg	0.549	0.502
Dissimilar	Bookshelf	0.239	0.226
	Bottle	0.174	0.184
	Clock	0.257	0.281
	Microwave	0.251	0.242
	Pianoforte	0.168	0.188
	Telephone	0.252	0.219
	Avg	0.223	0.224

Table . Testing Loss results for Point-Cloud Completion (training on all 7 known classes jointly using Vanilla Loss and Matching Loss). Testing Loss: mean Chamfer Distance only on the missing part.

Outputs for Matching Loss model



Point-cloud completion: Pros, cons and improvements

PROS

- ✓ Good at reconstructing general shape of the object
- ✓ Good at “transfer” local features learned for an object to reconstruct another one
- ✓ Small embedding size

CONS

- ✗ Bad at reconstructing finer details
- ✗ Bad at preserving symmetry

IMPROVEMENTS



Adjust the loss function to preserve symmetry and better reconstruct finer details

Thanks for your attention!




Image 'Smile - Statistical Machine Intelligence and Learning Engine (haifengl.github.io)'

Luca Barco s276072@studenti.polito.it

Stefano Bergia s276124@studenti.polito.it

Daniela De Angelis s277493@studenti.polito.it

Politecnico di Torino
Machine Learning and Artificial Intelligence
A.Y. 2020/2021



Image: polito.it